Minsik Choi (최민식)

Integrated B.S./M.S. Student, Korea University | Efficient LLMs, Model Merging & Transformer Optimization
mszzang2002@korea.ac.kr | +82 10-8700-5813 | Seoul, South Korea

Research Interests

- Model merging for parallel post-training of LLMs: combining experts trained separately via SFT or RLVR
- Understanding how post-training and adaptation reshape models, from spectral structure in weight space to attention behavior
- Efficient transformer inference: pruning and hardware-aware optimization for AI accelerators

Education

Korea University Sep. 2025 – Aug. 2027 (expected)
Integrated B.S./M.S. Program, Department of Computer Science and Engineering *Seoul, Korea*
- Advisor: Prof. Young Geun Kim

University of Tübingen Sep. 2023 – Feb. 2024
Exchange Student, Computer Science *Tübingen, Germany*

Korea University Mar. 2022 – Feb. 2026
B.S. in Computer Science and Engineering *Seoul, Korea*
- Admitted via the School of Interdisciplinary Studies · GPA: 4.2 / 4.5

Research Experience

FuriosaAI Sep. 2026 – Present
AI Research Engineering Intern, Algorithm Team *Seoul, Korea*

NAVER Cloud Sep. 2025 – Mar. 2026
Research & Engineering Intern *Seongnam, Korea*
- Contributed to a government-funded sovereign foundation model initiative (독자 파운데이션 모델 프로젝트), supporting the development of proprietary large language models.
- Designed and ran research on LLM training and evaluation pipelines, including instruction tuning and model merging.

Intelligent Computer Architecture & Systems Lab Feb. 2025 – Feb. 2026
Undergraduate Research Intern, Korea University (Advisor: Prof. Young Geun Kim) *Seoul, Korea*
- Hardware-aware transformer optimization and structured pruning for efficient on-device NLP inference.

Publications

*co-first author †corresponding author

Peer-Reviewed Conference Papers

[C1] Geeho Kim*, Minsik Choi*, Kyle Min, Young Geun Kim†, Bohyung Han†, "Merging RLVR-Trained Experts via Policy-Shift-Guided Spectral Alignment," *Proc. of the 2026 Conference on Neural Information Processing Systems (NeurIPS)*, 2026. **Main Conference**.

[C2] Minsik Choi*, Geewook Kim*, Young Geun Kim, "Text Capability Loss in Vision-Language Adaptation: An Attention-Sink Diagnosis," *Proc. of the 2026 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2026. **Main Conference**. poster

[C3] Minsik Choi*, Hyegang Son*, Changhoon Kim, Young Geun Kim, "Entropy Meets Importance: A Unified Head Importance–Entropy Score for Stable and Efficient Transformer Pruning," *Findings of EMNLP*, 2026. Extended version of [W1].

[C4] Minsik Choi*, Geewook Kim*†, "Decentralized Instruction Tuning: Conflict-Aware Splitting and Weight Merging," *Proc. of the 43rd International Conference on Machine Learning (ICML)*, 2026. poster

Peer-Reviewed Workshop Papers

[W1] Minsik Choi, Hyegang Son, Joohun Hyun, Seokmin Kim, Young Geun Kim, "Entropy Meets Importance: A Unified Head Importance–Entropy Score for Stable and Efficient Transformer Pruning," *NeurIPS 2025 Workshop*. poster

Technical Reports

[**T1**] NAVER Cloud HyperCLOVA X Team, "HyperCLOVA X 8B Omni," arXiv:2601.01792, 2026.

[**T2**] NAVER Cloud HyperCLOVA X Team, "HyperCLOVA X 32B Think," arXiv:2601.03286, 2026.

Tech Blog Posts

[**B1**] Minsik Choi, Kiuk Kim, "충돌은 나누고, 성능은 합치다: 효율적인 LLM Post-training 전략 MERIT," *NAVER CLOVA Tech Blog*, Jul. 2026. Covers [C4]. link

Academic Service

Reviewer, OPT 2026: Optimization for Machine Learning 2026
NeurIPS 2026 Workshop

Technical Skills

Research areas: LLM post-training (instruction tuning, RLVR / reinforcement fine-tuning); model merging and weight-space arithmetic for distributed post-training; transformer efficiency and structured / attention-head pruning; vision-language and multimodal models; efficient on-device inference.
Methods & analysis: spectral and geometric analysis of task interactions in weight space; head importance–entropy scoring; attention-sink diagnosis for vision-language adaptation; large-scale training and evaluation pipeline design.
Programming & frameworks: Python, PyTorch, Hugging Face Transformers; distributed / large-scale training (DeepSpeed, FSDP); CUDA; Git, Linux.
Languages: Korean (native), English (fluent).

Scholarships & Awards

Graduate Merit-Based Scholarship, Korea University Feb. 2026 – Aug. 2027
50% tuition, all semesters — awarded for academic excellence

Woonhae Scholarship Foundation Academic Scholarship 2025
Merit-based scholarship covering full tuition for one academic year